

## Driving - fork()

노트북: SW

만든 날짜: 2018-04-22 오후 7:33

수정한 날짜: 2018-04-29 오후 8:59

작성자: fstopdg@gmail.com

---

### How to find sys\_fork

1.  
In kernel/linux-4.4  
-> vi -t task\_struct (enter into 144)  
-> :cs find 0 fork (enter into 76 - kernel/fork.c)  
2.  
On vim,  
-> grep -rn "fork" | grep "SYSCALL"  
-> vi kernel/fork.c  
-> /SYSCALL\_DEFINE(fork)  
  
1791 SYSCALL\_DEFINE0(fork) -  
1792 {  
1793 #ifdef CONFIG\_MMU  
1794 return \_do\_fork(SIGCHLD, 0, 0, NULL, NULL, 0);  
\_do\_fork  
parameter : SIGCHLD, 0, 0, NULL, NULL, 0

### Enter into '\_do\_fork'

```
/* SIGCHLD, 0, 0, NULL, NULL, 0 */
long _do_fork(unsigned long clone_flags,
              unsigned long stack_start,
              unsigned long stack_size,
              int __user *parent_tidptr,
              int __user *child_tidptr,
              unsigned long tls)
{
    struct task_struct *p;
    int trace = 0;
    long nr;

    /*
     * Determine whether and which event to report to ptracer. When
     * called from kernel_thread or CLONE_UNTRACED is explicitly
     * requested, no event is reported; otherwise, report if the event
     * for the type of forking is enabled.
     */

    /*CLONE_UNTRACED = 0x00800000 */
    if (!(clone_flags & CLONE_UNTRACED)) {
        /*CLONE_VFORK = 0x00004000*/
        if (clone_flags & CLONE_VFORK)
            trace = PTRACE_EVENT_VFORK;
        /*CSIGNAL = 0x000000ff
        if clone_flags != SIGCHLD, trace = PTRACE_EVENT_CLONE = 3
        if clone_flag = SIGCHLD, trace = PTRACE_EVENT_FORK = 1 */
        else if ((clone_flags & CSIGNAL) != SIGCHLD)
            trace = PTRACE_EVENT_CLONE;
        else
            trace = PTRACE_EVENT_FORK;

        if (likely(!ptrace_event_enabled(current, trace)))
```

```

    trace = 0;
}

p = copy_process(clone_flags, stack_start, stack_size,
    child_tidptr, NULL, trace, tls);
/*
 * Do this prior waking up the new thread - the thread pointer
 * might get invalid after that point, if the thread exits quickly.
 */
if (!IS_ERR(p)) {
    struct completion vfork;
    struct pid *pid;

    trace_sched_process_fork(current, p);

    pid = get_task_pid(p, PIDTYPE_PID);
    nr = pid_vnr(pid);

    if (clone_flags & CLONE_PARENT_SETTID)
        put_user(nr, parent_tidptr);

    if (clone_flags & CLONE_VFORK) {
        p->vfork_done = &vfork;
        init_completion(&vfork);
        get_task_struct(p);
    }

    wake_up_new_task(p);

    /* forking complete and child started to run, tell ptracer */
    if (unlikely(trace))
        ptrace_event_pid(trace, pid);

    if (clone_flags & CLONE_VFORK) {
        if (!wait_for_vfork_done(p, &vfork))
            ptrace_event_pid(PTRACE_EVENT_VFORK_DONE, pid);
    }

    put_pid(pid);
} else {
    nr = PTR_ERR(p);
}
return nr;
}

```

There are 3 parts.

First, Above `p = copy_process(clone_flags, stack_start, stack_size, child_tidptr, NULL, trace, tls);`

Second, `p = copy_process(clone_flags, stack_start, stack_size, child_tidptr, NULL, trace, tls);`

Finally, Under `p = copy_process(clone_flags, stack_start, stack_size, child_tidptr, NULL, trace, tls);`

-> First section,

```

/* SIGCHLD, 0, 0, NULL, NULL, 0 */
long _do_fork(unsigned long clone_flags,
    unsigned long stack_start,
    unsigned long stack_size,
    int __user *parent_tidptr,
    int __user *child_tidptr,
    unsigned long tls)
{
    struct task_struct *p;

```

```

int trace = 0;
long nr;

/*
 * Determine whether and which event to report to ptracer. When
 * called from kernel_thread or CLONE_UNTRACED is explicitly
 * requested, no event is reported; otherwise, report if the event
 * for the type of forking is enabled.
 */

/*CLONE_UNTRACED = 0x00800000 */
if (!(clone_flags & CLONE_UNTRACED)) {
    /*CLONE_VFORK = 0x00004000*/
    if (clone_flags & CLONE_VFORK)
        trace = PTRACE_EVENT_VFORK;
    /*CSIGNAL = 0x000000ff
    if clone_flags != SIGCHLD, trace = PTRACE_EVENT_CLONE = 3
    if clone_flag = SIGCHLD, trace = PTRACE_EVENT_FORK = 1 */
    else if ((clone_flags & CSIGNAL) != SIGCHLD)
        trace = PTRACE_EVENT_CLONE;
    else
        trace = PTRACE_EVENT_FORK;

    if (likely(!ptrace_event_enabled(current, trace)))
        trace = 0;
}

```

$SIGCHLD = 17 \sim 20 = 10001 \sim 10100 = 0x11 \sim 0x14$

if -> else

trace = PTRACE\_EVENT\_FORK = 1

if(likely(!ptrace\_event\_enabled(current, trace)))

likely -> trace = 0 (iot depends on value of ptrace)

Anyway, trace = 0

-> Second section

```
p = copy_process(clone_flags, stack_start, stack_size, child_tidptr, NULL, trace, tls);
```

parameter : SIGCHLD, 0, 0, NULL, NULL, 0, 0

retval = security\_task\_create(clone\_flags)